

# EECS 545: Machine Learning

## Lecture 7. Kernel methods

Honglak Lee

1/31/2011



# Outline


- Recap: Kernel methods
- Support Vector Machines
- Next lecture: Gaussian Processes

# Kernel regression

# Radial Basis Functions

- Basis functions can be chosen that depend only on distance from selected centers:

$$\phi_j(\mathbf{x}) = h(\underbrace{\|\mathbf{x} - \mu_j\|}_{\text{distance}})$$



- A function  $f(x)$  can be approximated as a linear combination of the basis functions

$$\underline{f(\mathbf{x})} = \sum_{n=1}^N \underbrace{w_n}_{\text{weight}} h(\|\mathbf{x} - \mu_n\|)$$

- With a basis function at each training data point, the approximation is exact on the training data.

# Kernel Regression



- Using radial basis functions around the training data points, predict a value  $y(\mathbf{x})$  as the average of target values  $t_n$ , weighted by similarities  $k(\mathbf{x}, \mathbf{x}_n)$ :

$$y(\mathbf{x}) = \sum_{n=1}^N \underbrace{k(\mathbf{x}, \mathbf{x}_n)}_{\propto \exp\left[-\frac{\|\mathbf{x} - \mathbf{x}_n\|^2}{2\sigma^2}\right]} t_n$$

# Kernel Normalization

- The weighted average approach assumes

$$\sum_{n=1}^N k(\mathbf{x}, \mathbf{x}_n) = 1 \quad , \quad \forall \mathbf{x}$$



# Narayada-Watson model

- From kernel density estimation:

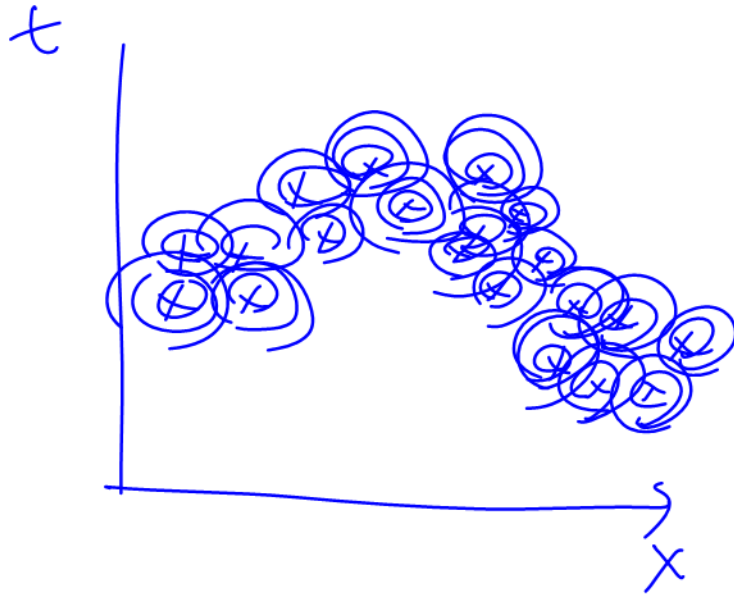
$$p(\mathbf{x}, t) = \frac{1}{N} \sum_{n=1}^N f(\mathbf{x} - \mathbf{x}_n, t - t_n)$$

data feature


output

Joint density function centered at  $x_n, t_n$

- where  $f(x, t)$  is the component density function and there is one such component centred on each data point
- We now find an expression for the regression function  $y(x)$ , corresponding to the conditional average of the target variable conditioned on the input variable



$$f(x) = \sum_n f(x - x_n, t - t_n)$$





# Narayada-Watson model

$$p(x, t)$$

$$E[t|x]$$

$$y(\mathbf{x}) = \mathbb{E}[t|\mathbf{x}] = \int_{-\infty}^{\infty} t p(t|\mathbf{x}) dt$$

$$= \frac{\int t p(\mathbf{x}, t) dt}{\int p(\mathbf{x}, t) dt}$$

$$= \frac{\sum_n \int t f(\mathbf{x} - \mathbf{x}_n, t - t_n) dt}{\sum_m \int f(\mathbf{x} - \mathbf{x}_m, t - t_m) dt}$$

$$p(t|x) = \frac{p(t, x)}{\int p(t, x) dt}$$

$$t = (t - t_n) + t_n$$

(6.43)

$$= g(\mathbf{x} - \mathbf{x}_m)$$

We now assume for simplicity that the component density functions have zero mean so that

$$\int_{-\infty}^{\infty} f(\mathbf{x}, t) t dt = 0 \quad (6.44)$$

for all values of  $\mathbf{x}$ . Using a simple change of variable, we then obtain

$$y(\mathbf{x}) = \frac{\sum_n g(\mathbf{x} - \mathbf{x}_n) t_n}{\sum_m g(\mathbf{x} - \mathbf{x}_m)} = k(\mathbf{x}, \mathbf{x}_n)$$

$$= \sum_n k(\mathbf{x}, \mathbf{x}_n) t_n \quad (6.45)$$

$$\sum_n k(\mathbf{x}, \mathbf{x}_n) = 1$$

# Narayada-Watson model

- Prediction function:

$$\begin{aligned}y(\mathbf{x}) &= \frac{\sum_n g(\mathbf{x} - \mathbf{x}_n) t_n}{\sum_m g(\mathbf{x} - \mathbf{x}_m)} \\&= \sum_n k(\mathbf{x}, \mathbf{x}_n) t_n\end{aligned}$$

– where

$$k(\mathbf{x}, \mathbf{x}_n) = \frac{g(\mathbf{x} - \mathbf{x}_n)}{\sum_m g(\mathbf{x} - \mathbf{x}_m)}$$


$$g(\mathbf{x}) = \int_{-\infty}^{\infty} f(\mathbf{x}, t) dt.$$

# Narayada-Watson model

- This model is also known as kernel regression.
- For a localized kernel function, it has the property of giving more weight to data points that are close to  $x$

# Kernel Regression Example

- On the familiar sinusoidal data set:



# Support Vector Machines

# Classification

- Consider a two-class classification problem:
  - Positive:  $t = +1$
  - Negative:  $t = -1$
- Train a linear model over the feature vector:





$$y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + b$$

- Train with input vectors  $\mathbf{x} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ 
  - and corresponding target values  $\mathbf{t} = \{t_1, \dots, t_N\}$ .
  - $y(\mathbf{x}) > 0 \Rightarrow t = +1$  and  $y(\mathbf{x}) < 0 \Rightarrow t = -1$
  - That is:  $t_n y(\mathbf{x}_n) > 0$ .

label  $\uparrow$  prediction for  $\mathbf{x}_n$

# Discriminant Function


- It can be arbitrary functions of  $\mathbf{x}$ , such as:



# Distance from Decision Surface

$$y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + b$$

- $w$  determines direction.
- $b$  determines offset.






# Linear Discriminant Function

- $g(\mathbf{x})$  is a linear function:

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

- A hyper-plane in the feature space
- (Unit-length) normal vector of the hyper-plane:


$$\mathbf{n} = \frac{\mathbf{w}}{\|\mathbf{w}\|}$$



# Linear Discriminant Function

- How would you classify these points using a linear discriminant function in order to minimize the error rate?


■ Infinite number of answers!



# Linear Discriminant Function

- How would you classify these points using a linear discriminant function in order to minimize the error rate?


■ Infinite number of answers!



# Linear Discriminant Function


- How would you classify these points using a linear discriminant function in order to minimize the error rate?

■ Infinite number of answers!



# Linear Discriminant Function

- How would you classify these points using a linear discriminant function in order to minimize the error rate?
- Infinite number of answers!
- Which one is the best?




# Large Margin Linear Classifier

$$\text{margin} = \min_{x^{(i)}} t^{(i)} (w^T x^{(i)} + b)$$



● denotes +1  
○ denotes -1

- The linear discriminant function (classifier) with the maximum **margin** is the best
- Margin is defined as the width that the boundary could be increased by before hitting a data point
- Why it is the best?
  - Robust to outliers and thus strong generalization ability




# Maximum Margin

- The margin is the minimum distance of an example from the decision surface.
- Determine  $w$  and  $b$  to maximize the margin.



# Support Vectors

- The *support vectors* are the points closest to the decision surface  $y(x) = 0$ .
- Set  $w$  so that  $t_n y(x_n) = 1$  for support vectors.
- Only the support vectors determine the decision surface.



$$t_n y(x_n) = \text{margin} \\ \geq 1, \forall n$$



# Constraints for Optimization

- Set  $w$  and  $b$  so that, for support vectors:

$$t_n(\mathbf{w}^T \phi(\mathbf{x}_n) + b) = 1$$

- Then *every* data point must satisfy:

$$t_n(\mathbf{w}^T \phi(\mathbf{x}_n) + b) \geq 1 \text{ for } n = 1, \dots, N$$

- It will turn out that only support vectors are active constraints.

# Large Margin Linear Classifier

● denotes +1  
○ denotes -1

- Given a set of data points:  
 $\{(\mathbf{x}_i, y_i)\}, i = 1, 2, \dots, n$ , where


For  $y_i = +1$ ,  $\mathbf{w}^T \mathbf{x}_i + b > 0$

For  $y_i = -1$ ,  $\mathbf{w}^T \mathbf{x}_i + b < 0$

- With a scale transformation on both  $\mathbf{w}$  and  $b$ , the above is equivalent to

For  $y_i = +1$ ,  $\mathbf{w}^T \mathbf{x}_i + b \geq 1$

For  $y_i = -1$ ,  $\mathbf{w}^T \mathbf{x}_i + b \leq -1$



# Large Margin Linear Classifier

● denotes +1  
○ denotes -1

- We know that

$$\mathbf{w}^T \mathbf{x}^+ + b = 1$$

$$\mathbf{w}^T \mathbf{x}^- + b = -1$$


$$\mathbf{w}^T (\mathbf{x}^+ - \mathbf{x}^-) = 2$$

- The margin width is:

$$M = (\mathbf{x}^+ - \mathbf{x}^-) \cdot \mathbf{n}$$

$$= (\mathbf{x}^+ - \mathbf{x}^-) \cdot \frac{\mathbf{w}}{\|\mathbf{w}\|} = \frac{2}{\|\mathbf{w}\|}$$

$$\hat{\mathbf{n}} = \frac{\mathbf{w}^T}{\|\mathbf{w}\|} (\mathbf{x}^+ - \mathbf{x}^-) = \frac{2}{\|\mathbf{w}\|}$$



# Large Margin Linear Classifier

● denotes +1  
○ denotes -1

- Formulation:

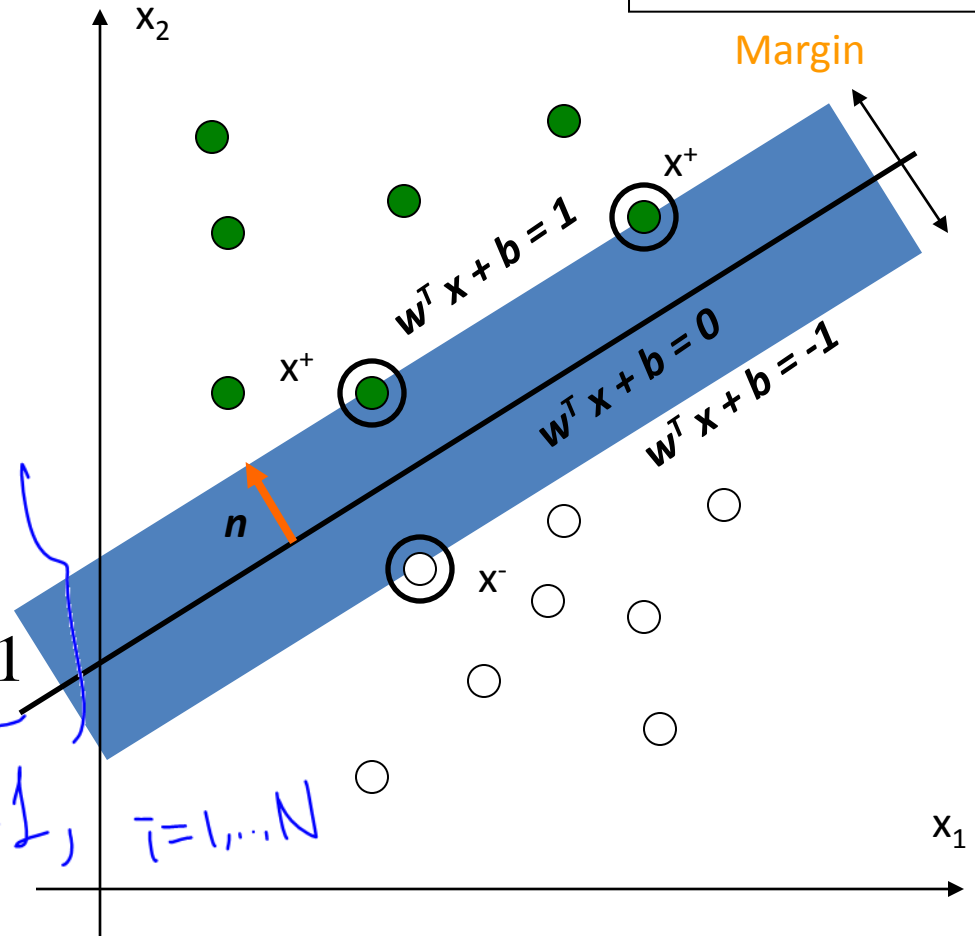
$$\text{maximize } \frac{2}{\|\mathbf{w}\|}$$

such that

For  $y_i = +1$ ,  $\mathbf{w}^T \mathbf{x}_i + b \geq 1$

For  $y_i = -1$ ,  $\mathbf{w}^T \mathbf{x}_i + b \leq -1$

$$t_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1, \quad i=1, \dots, N$$



# Maximize the Margin

- Distance to decision surface is  $y(x_n)/||\mathbf{w}||$
- To maximize the margin, maximize  $||\mathbf{w}||^{-1}$
- This is the same as minimizing  $||\mathbf{w}||^2$
- Use Lagrange multipliers to enforce constraints while optimizing

minimize

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} ||\mathbf{w}||^2 - \sum_{n=1}^N a_n \left\{ t_n (\mathbf{w}^T \phi(\mathbf{x}_n) + b) - 1 \right\}$$

Handwritten notes and constraints:

- $a_n \geq 0$  (written above the sum)
- $t_n (\mathbf{w}^T \phi(\mathbf{x}_n) + b) - 1 \geq 0$  (written inside the sum term)
- $\sum_{n=1}^N a_n \geq 0$  (written below the sum)
- $\frac{1}{2} ||\mathbf{w}||^2 \leq \frac{1}{2} ||\mathbf{w}||^2$  (circled at the bottom right)

$$\min_x f(x)$$

$$\text{s.t. } g_i^*(x) \leq 0$$

$$i = 1, \dots, N$$

$$j=1, \dots, M$$

$f$ : Convex

$g$ : convex

$h$ : affine

## 1. Lagrangian

Lagrangian

$$\mathcal{L}(x, \mu, \lambda) = f(x) + \sum_{i=1}^N \lambda_i g_i(x) + \sum_{j=1}^M \mu_j h_j(x)$$

Convex

2. Solve Lagrange dual.

$$\max_{\mu, \lambda} \min_x \mathcal{L}(x, \mu, \lambda)$$

$$\text{s.t. } \lambda_i \geq 0, i=1, \dots, N$$

afm.

$ax+b$

$$D(\mu, \lambda)$$

$$x - x_0 \approx \nabla f(x)$$

- $$f(x) = f(x_0) + \nabla f(x)^T (x - x_0)$$



# Lagrange Multipliers

- Since the gradients are parallel, there must exist a parameter (the Lagrange multiplier)

$$\nabla f + \lambda \nabla g = 0$$

$$g(\mathbf{x}) = 0$$

- Then we define the Lagrangian function

$$\underline{L(\mathbf{x}, \lambda) \equiv f(\mathbf{x}) + \lambda g(\mathbf{x})}$$

- to optimize:

$$\nabla_{\mathbf{x}} L(\mathbf{x}, \lambda)$$


$$\nabla_{\mathbf{x}} L = 0 \text{ implies } \nabla f + \lambda \nabla g = 0$$

$$\underline{\partial L / \partial \lambda} = 0 \text{ implies } \boxed{g(\mathbf{x}) = 0}$$



# Lagrange Multipliers

- Suppose we have an inequality constraint  $g(\mathbf{x}) \geq 0$
- If boundary optimum  $\mathbf{x}_A$  then gradient of  $f$  is outward, and  $\lambda > 0$
- If internal optimum  $\mathbf{x}_B$  then  $\lambda = 0$



$\begin{cases} \text{if } \mathbf{x}^* \text{ on surface } g(\mathbf{x})=0 \\ \text{if inside, } \end{cases} \quad \lambda \nabla g(\mathbf{x}) + \nabla f(\mathbf{x}) = 0$

# Lagrange Multipliers

- Combining these cases gives us the *Karush-Kuhn-Tucker (KKT) conditions* when maximizing  $f(x)$  subject to an inequality constraint.

$$g(\mathbf{x}) \geq 0$$

$$\lambda \geq 0$$

$$\boxed{\lambda g(\mathbf{x}) = 0}$$

$$\rightarrow \begin{cases} \lambda = 0 \\ \text{or } \lambda > 0, \text{ then } g(x) = 0 \end{cases}$$

# Maximize the Margin

- Distance to decision surface is  $y(x_n)/||\mathbf{w}||$
- To maximize the margin, maximize  $||\mathbf{w}||^{-1}$
- This is the same as minimizing  $||\mathbf{w}||^2$
- Use Lagrange multipliers to enforce constraints while optimizing

Lagrangian

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} ||\mathbf{w}||^2 - \sum_{n=1}^N a_n \{ t_n (\mathbf{w}^T \phi(\mathbf{x}_n) + b) - 1 \}$$

$\nabla_{\mathbf{w}} L = \mathbf{w}$

$a_n \geq 0$

$\sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n) \geq 0$

$\sum a_n t_n \phi(\mathbf{x}_n)$

# Maximize the Margin

- Set the derivatives of  $L(w, b, a)$  to zero, to get

$$\nabla_w L(w, b, a)$$

$$\underline{\mathbf{w}} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

$$\nabla_b L(w, b, a)$$

$$\underline{0 = \sum_{n=1}^N a_n t_n} //$$

- Substitute in, to eliminate  $w$  and  $b$ ,

$$\text{Lagrange dual}$$

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m \underbrace{\phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m)}_{//}$$

$$// \quad \langle \mathbf{x}_n, \mathbf{x}_m \rangle$$

# Dual Representation (with kernel)

- Define a kernel  $k(\mathbf{x}_n, \mathbf{x}_m) = \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m)$

- This gives, to maximize

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

s.t.  $a_n \geq 0$

- Once we have  $\mathbf{a}$ , we don't need  $\mathbf{w}$ . Predict new values using

$$y(\mathbf{x}) = \underbrace{\mathbf{w}^T \phi(\mathbf{x})}_{\sum_n a_n t_n \phi(\mathbf{x}_n)} + b = \sum_{n=1}^N a_n t_n k(\mathbf{x}, \mathbf{x}_n) + b$$

# Recovering b

- For any support vector  $x_n$ :  $t_n y(x_n) = 1$
- Replacing with  $y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + b = \sum_{n=1}^N a_n t_n k(\mathbf{x}, \mathbf{x}_n) + b$

$$t_n \left[ t_n \left( \sum_{m \in \mathcal{S}} a_m t_m k(\mathbf{x}_n, \mathbf{x}_m) + b \right) \right] = 1$$

$\uparrow$   
 (index) set of support vectors

$t_n^2 = 1$

- Multiply  $t_n$ , and sum over n:

$$b = \frac{1}{N_{\mathcal{S}}} \sum_{n \in \mathcal{S}} \left( t_n - \sum_{m \in \mathcal{S}} a_m t_m k(\mathbf{x}_n, \mathbf{x}_m) \right)$$

# Support Vectors

- The KKT conditions are:

$$a_n \geq 0$$

$$t_n y(\mathbf{x}_n) - 1 \geq 0$$

$$a_n \{t_n y(\mathbf{x}_n) - 1\} = 0$$

- Which means, either  $a_n=0$  or  $t_n y(\mathbf{x}_n)=1$ .
- That is, only the support vectors matter!
  - To predict  $y(\mathbf{x})$ , sum only over support vectors


# Support Vector Machines

- Hard SVM requires separable sets

$$t_n y(\mathbf{x}_n) - 1 \geq 0$$

- Soft SVM introduces *slack variables* for each data point



$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n$$





# Soft SVM

- A little slack can give much better margin.



# Soft SVM

- Maximize the margin, and also penalize for the slack variables

$$C \sum_{n=1}^N \xi_n + \frac{1}{2} \|\mathbf{w}\|^2$$

- The support vectors are now those with

$$t_n y(\mathbf{x}_n) = 1 - \xi_n$$

# Formulation of soft-margin SVM

- Primal form
- Minimize (w.r.t.  $w$  and  $\xi_n$ 's)

$$C \sum_{n=1}^N \xi_n + \frac{1}{2} \|\mathbf{w}\|^2$$

Subject to  $t_n y(\mathbf{x}_n) \geq 1 - \xi_n, \forall n$

$$\xi_n \geq 0, \forall n$$

# Dual formulation of soft-margin SVM

- Lagrangian

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

– Where  $a_n \geq 0$ ,  $\mu_n \geq 0$ ,  $\xi_n \geq 0, \forall n$

- KKT conditions for the constraints

$$a_n \geq 0$$

$$t_n y(\mathbf{x}_n) - 1 + \xi_n \geq 0$$

$$a_n (t_n y(\mathbf{x}_n) - 1 + \xi_n) = 0$$

$$\mu_n \geq 0$$

$$\xi_n \geq 0$$

$$\mu_n \xi_n = 0$$

# Dual formulation of soft-margin SVM

- Taking derivatives

$$\frac{\partial L}{\partial \mathbf{w}} = 0 \quad \Rightarrow \quad \mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

$$\frac{\partial L}{\partial b} = 0 \quad \Rightarrow \quad \sum_{n=1}^N a_n t_n = 0$$

$$\frac{\partial L}{\partial \xi_n} = 0 \quad \Rightarrow \quad a_n = C - \mu_n.$$

# Dual formulation of soft-margin SVM

- Lagrange dual

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

subject to  $0 \leq a_n \leq C$

$$\sum_{n=1}^N a_n t_n = 0$$

- Solve quadratic problem (convex optimization)

# Support Vector Machine: Algorithm

- 1. Choose a kernel function
- 2. Choose a value for  $C$
- 3. Solve the quadratic programming problem  
(many software packages available)
- 4. Construct the discriminant function from the support vectors

# Some Issues

- Choice of kernel
  - Gaussian or polynomial kernel is default
  - if ineffective, more elaborate kernels are needed
  - domain experts can give assistance in formulating appropriate similarity measures
- Choice of kernel parameters
  - e.g.  $\sigma$  in Gaussian kernel
  - $\sigma$  is the distance between closest points with different classifications
  - In the absence of reliable criteria, applications rely on the use of a validation set or cross-validation to set such parameters.



# Summary: Support Vector Machine

- 1. Large Margin Classifier
  - Better generalization ability & less over-fitting
- 2. The Kernel Trick
  - Map data points to higher dimensional space in order to make them linearly separable.
  - Since only dot product is used, we do not need to represent the mapping explicitly.

# Additional Resource

- <http://www.kernel-machines.org/>

# SVM Implementation

- LIBSVM
  - <http://www.csie.ntu.edu.tw/~cjlin/libsvm/>
  - One of the most popular generic SVM solver (supports nonlinear kernels)
- Liblinear
  - <http://www.csie.ntu.edu.tw/~cjlin/liblinear/>
  - One of the fastest linear SVM solver
- SVMlight
  - [http://www.cs.cornell.edu/people/tj/svm\\_light/](http://www.cs.cornell.edu/people/tj/svm_light/)
  - Structured outputs, various objective measure (e.g., F1, ROC area), Ranking, etc.

# SVM demo code

- <http://www.mathworks.com/matlabcentral/fileexchange/28302-svm-demo>
- <http://www.alivelearn.net/?p=912>