

EECS 545: Machine Learning

Lecture 10. Regularization and model selection

Honglak Lee

2/9/2011




Outline

- ML, MAP, and Bayesian
 - Maximum Likelihood
 - MAP
 - Bayesian
- Bias-Variance Tradeoff
- Model selection
 - Cross validation
- Advice on applying Machine Learning

ML vs. MAP vs. Bayesian (summary)

- Maximum Likelihood
 - Objective: Log-likelihood
 $\log P(D|w)$
 - Example: linear regression (w/o regularization)
- MAP (Maximum a Posteriori)
 - Objective: Log-likelihood + Log-Prior
 $\log P(D|w) + \log P(w)$
 - Example: Regularized linear regression
- Bayesian
 - Objective: Estimate $P(w|D) \propto P(D|w)P(w)$
 - Goal is not a point estimate!! (Unlike ML or MAP)
 - Example: Bayesian linear regression





Polynomial Curve Fitting




$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

Maximum Likelihood (in Linear Regression)

- Choosing the right complexity is important
 - Watch out for underfitting/overfitting



Underfitting vs. overfitting





Root-Mean-Square (RMS) Error: $E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$

How can we avoid overfitting?

- More training data
 - Always helps
- Regularization
- Bayesian



More training data

- Even complicated models can benefit (avoid overfitting) by having large amount of data
 - Example: 9th order polynomial ($M=9$)



Regularization (for Linear Regression)

- Regularization can implicitly control the complexity of models
 - Example: 9th order polynomial (M=9)
 - Choosing right level of regularization is important




$$E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

Data term + Regularization term

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

Bayesian (Linear Regression)



Bayesian still assumes some hyperparameters, but overfitting is generally much less than maximum-likelihood

ML vs. MAP vs. Bayesian (summary)

- Maximum Likelihood
 - Objective: Log-likelihood
 $\log P(D|w)$
 - Example: linear regression (w/o regularization)
- MAP (Maximum a Posteriori)
 - Objective: Log-likelihood + Log-Prior
 $\log P(D|w) + \log P(w)$
 - Example: Regularized linear regression
- Bayesian
 - Objective: Estimate $P(w|D) \propto P(D|w)P(w)$
 - Goal is not a point estimate!! (Unlike ML or MAP)
 - Example: Bayesian linear regression





Variance-Bias Tradeoff

The Bias-Variance Decomposition

- Setting:
 - Given a distribution of $P(x,y)$
 - Sample training data
$$D_{train} = \{(x_n, y_n): n = 1, \dots, N\} \sim P(x, y)$$
 - Train a learning algorithm on D
- Depending on samples, learning algorithm can still give different results (ML, MAP, etc.)
- Goal: We want to learn a model with
 - Small bias (i.e., how well a model fits the data?)
 - Small variance (i.e., how stable a model is wrt data samples?)



The Bias-Variance Decomposition

- Example: samples from the sinusoidal function $y=\sin(x)$





The Bias-Variance Decomposition

- Example: 25 data sets from the sinusoidal, varying the degree of regularization λ .





The Bias-Variance Decomposition

- Example: 25 data sets from the sinusoidal, varying the degree of regularization λ .



The Bias-Variance Decomposition

- Example: 25 data sets from the sinusoidal, varying the degree of regularization λ .



The Squared Loss Function

Let's assume that we sample (\mathbf{x}, t) according to a probability distribution $p(\mathbf{x}, t)$. Then expectation of squared loss can be written as:

$$\mathbb{E}[L] = \iint \{y(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) d\mathbf{x} dt$$

Decompose:

$$\begin{aligned} \{y(\mathbf{x}) - t\}^2 &= \{y(\mathbf{x}) - \mathbb{E}[t|\mathbf{x}] + \mathbb{E}[t|\mathbf{x}] - t\}^2 \\ &= \{y(\mathbf{x}) - \mathbb{E}[t|\mathbf{x}]\}^2 + 2\{y(\mathbf{x}) - \mathbb{E}[t|\mathbf{x}]\}\{\mathbb{E}[t|\mathbf{x}] - t\} + \{\mathbb{E}[t|\mathbf{x}] - t\}^2 \end{aligned}$$

We finally get:

$$\mathbb{E}[L] = \int \{y(\mathbf{x}) - \mathbb{E}[t|\mathbf{x}]\}^2 p(\mathbf{x}) d\mathbf{x} + \int \text{var}[t|\mathbf{x}] p(\mathbf{x}) d\mathbf{x}$$

The Bias-Variance Decomposition (1)

- Recall the *expected squared loss*,

$$\mathbb{E}[L] = \int \{y(\mathbf{x}) - h(\mathbf{x})\}^2 p(\mathbf{x}) d\mathbf{x} + \underbrace{\int \int \{h(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) d\mathbf{x} dt}_{\text{noise}}$$

- where

$$h(\mathbf{x}) = \mathbb{E}[t|\mathbf{x}] = \int t p(t|\mathbf{x}) dt.$$

- The second term of $\mathbb{E}[L]$ corresponds to the noise inherent in the random variable t .
- What about the first term?

The Bias-Variance Decomposition (2)

- Suppose we were given multiple data sets, each of size N . Any particular data set, \mathcal{D} , will give a particular function $y(\mathbf{x}; \mathcal{D})$. We then have

$$\begin{aligned} & \{y(\mathbf{x}; \mathcal{D}) - h(\mathbf{x})\}^2 \\ &= \{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] + \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2 \\ &= \{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}^2 + \{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2 \\ &\quad + 2\{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}\{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}. \end{aligned}$$

The Bias-Variance Decomposition (3)

- Taking the expectation over \mathcal{D} yields

$$\begin{aligned} \mathbb{E}_{\mathcal{D}} [\{y(\mathbf{x}; \mathcal{D}) - h(\mathbf{x})\}^2] \\ = \underbrace{\{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2}_{(\text{bias})^2} + \underbrace{\mathbb{E}_{\mathcal{D}} [\{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}^2]}_{\text{variance}}. \end{aligned}$$

The Bias-Variance Decomposition (4)

- Thus we can write

$$\text{expected loss} = (\text{bias})^2 + \text{variance} + \text{noise}$$

- where

$$(\text{bias})^2 = \int \{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2 p(\mathbf{x}) \, d\mathbf{x}$$


$$\text{variance} = \int \mathbb{E}_{\mathcal{D}} [\{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}^2] p(\mathbf{x}) \, d\mathbf{x}$$

$$\text{noise} = \iint \{h(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) \, d\mathbf{x} \, dt$$

The Bias-Variance Trade-off

From these plots, we note that

- An over-regularized model (large λ) will have a high bias,
- While an under-regularized model (small λ) will have a high variance.



Model Selection

Choosing right models

- For polynomial curve fitting, which value of M should we choose?
- For SVM, which hyperparameter values should we choose?
 - Linear SVM: C (penalty for margin violation)
 - RBF kernel SVM: C , kernel width
 - Polynomial kernel SVM: C , degree of polynomial
- Generally, given a set of models, $M = \{M_1, M_2, \dots, M_d\}$, how can we choose optimal M_i ?
 - Model:
 - Class (or set) of hypothesis: learning algorithm, hyperparameters, etc.
 - Fixed during training
 - Parameter:
 - Aka, hypothesis: (w values for logistic regression/SVM/linear regression)
 - Can be trained based on data.

Simple Idea (that doesn't work)

- Given Data D
- Train each model M_i on D , to get some hypothesis (model) h_i
- Pick the hypothesis with the smallest training error

Cross validation

- Hold-out cross validation (aka simple cross validation)
 - 1. Randomly split D into D_{train} (say, 70% of the data) and D_{CV} (the remaining 30%).
 - Here, D_{CV} is called the hold-out cross validation set.
 - 2. Train each model M_i on D_{train} only, to get some hypothesis h_i .
 - 3. Select and output the hypothesis h_i that had the smallest error on the hold out cross validation set.
- Disadvantage:
 - Waste 30% of the data!

K fold Cross validation

- 1. Randomly split D into k disjoint subsets of N/k training examples each.
 - Lets call these subsets D_1, \dots, D_k .
- 2. For each model M_i , we evaluate it as follows:
 - For $j = 1, \dots, k$
 - Train the model M_i on $D_1 \cup \dots \cup D_{j-1} \cup D_{j+1} \cup \dots \cup D_k$ (i.e., train on all the data except D_j) to get some hypothesis h_{ij} .
 - Test the hypothesis h_{ij} on D_j , to get $\epsilon_{D_j}(h_{ij})$.
 - The estimated generalization error of model M_i is then calculated as the average of the $\epsilon_{D_j}(h_{ij})$'s (averaged over j).
- 3. Pick the model M_i with the lowest estimated generalization error, and retrain that model on the entire training set S . The resulting hypothesis is then output as our final answer.

K fold Cross validation

- Popular choice of $k = 10, 5$
- Special case: $K=1$
 - Called, Leave-one-out cross validation (LOO CV)
 - Expensive, but wastes least amount of training data for cross validation.